

Artificial Intelligence (CS 452/552)
Assignment 03 (60 points)
Due before class, Wednesday, 01 November 2017

Your goal for this assignment is to write a program that reads in an input file; this file will describe a Bayes Net in terms of

- a. A list of variables, each with their possible values.
- b. A list of variables, each with their parent variables in the BN.
- c. A list of variables, each with their CPTs.

As an example, consider the sample file `burglary.txt`:

```
Burglary T F
Earthquake T F
Alarm T F
JohnCalls T F
MaryCalls T F
# Parents
Alarm Burglary Earthquake
JohnCalls Alarm
MaryCalls Alarm
# Tables
Burglary
0.001
Earthquake
0.002
Alarm
T T 0.95
T F 0.94
F T 0.29
F F 0.001
JohnCalls
T 0.90
F 0.05
MaryCalls
T 0.70
F 0.01
```

This file describes the burglary BN from the text (p. 512). First, it shows each variable, with its possible values (T of F). Following the `# Parents` tag, we have each variable, followed by its

parents, and following the `# Tables` tag, we have the CPTs for each variable. Note that the rows of the CPT are ordered by the parents as given in the previous section of the input (if there are any), and contain numbers for the first $(d - 1)$ values of the variable (where d is its domain size). For example, each row of the CPT for the `Alarm` variable starts with a value for its first parent (`Burglary`), followed by a value for its second parent (`Earthquake`), followed by the probability $P(\text{Alarm} = T \mid \text{Parents})$.

Look at the various input files to make sure you understand all the details before you begin. In particular, *do not assume* that all variables are always boolean. The sample files also include `sprinklers.txt` (from p. 529, Figure 14.12 (a)), `books.txt` (from p. 539, Figure 14.17 (a)), and `data.txt` (from the BN exercises done in class earlier in the semester).

Your program will read in the file, using command line arguments, and then create a Bayes Net corresponding to that file. It will then allow the user to enter queries interactively, using standard input. A query will be in the form:

`<VARIABLE> | <EVIDENCE_1> = <VAL_1>, ..., <EVIDENCE_m> = <VAL_m>`

If there is no evidence, then the user will simply enter the name of a variable. Output will be in the form of a distribution over all the values of the query variable. When the user types `quit`, the program should terminate. A sample run would be as follows:

```
> java BayesNet burglary.txt

Loading file "burglary.txt".

Burglary
P(T) = 0.001, P(F) = 0.999

JohnCalls | Alarm = T
P(T) = 0.9, P(F) = 0.1

Burglary | JohnCalls = T, MaryCalls = T
P(T) = 0.284, P(F) = 0.716

quit
```

Note: if possible, your output should look exactly like this. It should be as close as possible, in any case. One thing to note is that the decimal values are formatted to have a maximum of 3 digits precision. The download for this assignment contains a number of sample runs.

Your program will use the enumerative query algorithm—outlined in the text (p. 525, Figure 14.9) and covered in lecture 16—to do its work.

You will hand in source-code for the program. The instructor will compile it and test it by giving it different input files. Two sample files are provided, and you can create your own for testing, but you need to use the same format as given above. Submit the work by email by the due-date and time (before class on the date given at the start of the assignment). Your email should include an archive containing the source code and instructions for using it.

I prefer that you code the solution in Java, but if you wish to use another language, let me know first. When doing your coding, follow these conventions:

1. Your program should run from the command-line, and needs to take a string as command-line input, as shown in the sample runs. The string will be the name of a text file containing a BN specification.
2. Provide a short **README** file with your code, explaining exactly what commands or procedures are needed to make it compile and run properly.
3. Follow basic software engineering principles; that is, your code should be clean and well-structured, with comments explaining each method or function, and the usual format conventions to make it readable.